
REPORT · 2026-06-30

Sergio Alberto Dominguez

AUTHOR

DATE

2026-06-30

AI Evaluation Specialist · Terminal-Bench Contributor · Technical Benchmark Author

San Luis, Argentina

Email: dominguez-sergioa@hotmail.com · Phone: +54 9 11 3135-3129

LinkedIn: <https://www.linkedin.com/in/sergio-alberto-dominguez-8a176546/> · GitHub: [\[tu-url-github.com\]](#)

Professional Summary

AI Evaluation Specialist with **15+ years** of hands-on technical depth, currently designing and validating complex benchmark tasks for AI agents at Revelo. Builds Terminal-Bench 2.0 tasks covering Software Engineering, Systems & Infrastructure, Security Engineering, Data Science, and Scientific Computing — including Linux hardening, MVCC snapshot isolation, BM25/MinHash ranking, LSM-tree compaction, and sparse conjugate gradient pipelines.

Combines **production-grade engineering background** (AWS, Azure, Docker, Kubernetes, Python, FastAPI, Terraform, CloudFormation) with rigorous **evidence-based review discipline**: writes oracle solutions, builds test harnesses, validates coverage, runs live Docker verifications, and produces detailed Q0–Q5 technical questionnaires with reproducible evidence. Familiar with SWE-Bench / Terminal-Bench style workflows and SARIF-based security audits.

Key Achievements

- Designed and validated **20+ AI benchmark task families** spanning Linux hardening, MVCC, LSM-trees, BM25/MinHash ranking, hermetic build graphs, sparse linear algebra, and TypeScript CLIs.
 - Achieved **Oracle PASS 35/35 tests, Coverage PASS 93.6%, Complexity PASS, Long-horizon PASS** on a Linux hardening benchmark task.
 - Completed **20 Vulnerability Audit tasks** for Supernova Full-Stack Vulnerability Audit v2, evaluating security annotations against SARIF reports, source code, and live Docker verification.
 - Audited **Infrastructure-as-Code submissions** across AWS CDK, Terraform, Pulumi, and CloudFormation with evidence-based verdicts.
 - Built and validated validators for **oracle correctness, coverage, complexity, and long-horizon behavior**.
-

Relevant Work Experience

Revelo — AI Training / Terminal Task Developer & Technical Evaluator

Freelance / Remote | 2026 - Present

- Designed, improved, and validated complex **Terminal-Bench 2.0 tasks** used to evaluate AI agents in terminal-driven development workflows.
- Created and reviewed hard technical scenarios covering **Software Engineering, Systems & Infrastructure, Security Engineering, Data Science, and Scientific Computing**.
- Built task specifications, oracle solutions, test suites, validation evidence, and feedback harnesses to measure whether AI agents can reason, code, execute scripts, and debug in terminal environments.
- Worked with **Python, JavaScript, TypeScript, Ruby, Bash, Docker, Linux CLI tooling, automated tests, and validation pipelines**.
- Developed hard benchmark tasks including **Linux hardening enforcement, MVCC snapshot isolation engine, hermetic build graph evaluator, BM25/MinHash document ranking, LSM-tree compaction, and sparse conjugate gradient pipelines**.
- Delivered validated benchmark evidence: **Oracle PASS 35/35, Coverage PASS 93.6%, Complexity PASS, Long-horizon PASS**.
- Audited Infrastructure-as-Code submissions for correctness, execution quality, test coverage, and evidence consistency across **AWS CDK, Terraform, Pulumi, and CloudFormation**.
- Reviewed AWS architecture patterns involving VPC, RDS PostgreSQL, S3, Lambda, SQS, EventBridge, Step Functions, API Gateway, AWS Backup, VPC Endpoints, and IAM policies.

Supernova Full-Stack Vulnerability Audit v2 — Security Evaluation

- Evaluated **20 pre-existing vulnerability audit tasks** against Q0–Q5 questionnaire criteria for AI Teacher submission.
- Performed **live Docker verification** of reported vulnerabilities in Flask 3 + React 19 + PostgreSQL 15 applications (EventHive).
- Analyzed **SARIF reports**, source code, and team annotations across 5 teams (jaylbreaK, lionxa, pr1smcode, purcl, seclab).
- Produced validated questionnaire responses with technical accuracy review, evidence-based verdicts, and JSON snippet references.
- Verified vulnerabilities including **mass assignment, broken authentication, SQL injection, XSS, and insecure deserialization** through actual HTTP requests.
- Delivered consistent evaluation framework with .status tracking, diagnosis contracts, and validation pipelines.

Charles Taylor InsureTech — Client Service / Cloud Infrastructure Engineer

San Luis, Argentina | Aug 2016 - Present (9 years)

- Implement cloud solution changes for multiple Charles Taylor clients in **production and UAT environments**.
- Operational support for **30+ environments** running **~190 servers** and **~120 virtual machines**.
- Provide **incident response, troubleshooting, technical documentation, and environment-specific guidance**.
- Work with **AWS CloudWatch, ServiceNow, Docker, Kubernetes, Linux, Windows Server, and cloud services**.
- Document runbooks and post-incident reports used by both internal teams and client stakeholders.

Earlier Technical Roles

- **Ser Beef S.A. — Help Desk L2 / Infrastructure**, San Luis, Argentina | Nov 2010 - Mar 2015
(Windows Server, AD, MySQL, MSSQL, networking)
 - **Dell — On-Site Technical Support**, San Luis, Argentina | Mar 2009 - Nov 2010
 - **Universidad de La Punta — Technical Support**, San Luis, Argentina | Feb 2007 - Mar 2009
-

Technical Skills

Benchmark & Evaluation: Terminal-Bench, SWE-Bench style workflows, oracle solution design, test harness construction, validator writing (oracle, coverage, complexity, long-horizon)

Security Review: SARIF report analysis, vulnerability triage, mass assignment / broken auth / SQL injection / XSS / insecure deserialization, live Docker verification

Programming: Python, TypeScript, JavaScript, Ruby, Bash, SQL

Cloud & Infra (for benchmark realism): AWS (CDK, Terraform, CloudFormation, IAM, VPC, RDS, Lambda, S3, SQS, EventBridge), Azure, Docker, Kubernetes, Linux

Backend / Data: FastAPI, SQLAlchemy, PostgreSQL, MySQL, Redis, ChromaDB, RAG pipelines

Validation Tooling: pytest, CLI validation, automated test design, evidence-based reviews, technical audits

Certifications & Training

Completed:

- AWS Certified Cloud Practitioner
- Microsoft Certified: Azure Fundamentals (AZ-900)

- Microsoft 365 Fundamentals (MS-900)
- Cloud Computing AWS (course)

In progress:

- AZ-104: Microsoft Azure Administrator

Additional training: DevOps Essentials, Azure Policy, Kubernetes for IT Administrators (LinkedIn Learning)

Education

National University of San Luis

Computer Engineering — In progress

Languages

- **Spanish:** Native
 - **English:** A2 (reading technical documentation; conversational proficiency in active development)
 - **Portuguese:** Basic
-

Why AI evaluation work fits me

- I write code that *fails* on purpose — benchmark tasks need to be hard enough that frontier agents get them wrong, and I know what hard looks like.
 - I'm comfortable across the stack — Linux hardening, storage engines, build systems, IaC, web apps, security. The variety is the job.
 - I treat every evaluation as an **evidence problem**: oracle, coverage, complexity, long-horizon, reproducibility. If the evidence isn't there, the verdict isn't either.
 - I produce written justifications a reviewer can audit. Q0–Q5 questionnaire work taught me to defend every claim.
-